Sovereign Sentinel Architecture V1.1

A Six-Axis Deterministic Framework for Frontier AI Safety
with Semantic-Referential Integrity Enforcement

Frank Bruno
AI Safety Auditor & Logic Architect
frank.bruno.oe@gmail.com
March 13, 2026

Abstract

We present the Sovereign Sentinel Architecture (SSA) V1.1, a proposed defense-in-depth control framework for frontier language model deployments operating across six distinct abstraction levels: mathematical, electronic, cryptographic, statistical-behavioral, semantic-referential, and institutional. The architecture addresses three documented failure modes: *Safety Amnesia* (erosion of safety constraints across extended context windows), *Stochastic Sabotage* (incremental assembly of prohibited knowledge through individually benign queries), and *Goal-Oriented Factual Inversion* (GOFI)—a non-adversarial failure mode in which a model correctly identifies factual ground truth in early turns, then generates output directly contradicting that ground truth when a persuasive goal frame is introduced in subsequent turns.

The GOFI failure mode was identified through empirical forensic audit—Scenario 5, March 2026—conducted against four deployed frontier models (Gemini, ChatGPT, Claude, and Copilot) across two party-inverted contract analysis scenarios in English and Spanish. All four models exhibited directional compliance failures. One model demonstrated correct first-turn identification of contractual risk, followed by a factually inverted second-turn deliverable whose inversion was directionally consistent with the prompted persuasive goal. This pattern is distinguished from generic analytical error by three criteria: directional consistency with the optimization target, correct prior processing of the contradicted fact, and professional formatting that renders the inversion undetectable without independent document review. Cryptographically anchored transcripts and all model deliverables from Scenario 5 are publicly archived for independent verification.

GOFI is not intercepted by any of the five preceding SSA axes. The sixth axis—the Deterministic Truth-Anchor with Fact-Claim Inversion Recognition (DTA-FCIR)—closes this gap through an immutable session-scoped Structured Fact Registry and an isolated formal Contradiction Engine operating at response granularity. The architecture has not been empirically validated at scale; all specifications presented are engineering targets proposed to meet defined safety invariants, and their formal correctness requires evaluation by researchers with the relevant technical expertise. Collaborative partnership with a Tier 1 AI safety team is sought for Phase 0 proof-of-concept implementation and adversarial evaluation.

Methodology and Disclosure: In alignment with emerging standards for transparent AI research, the author acknowledges the use of Large Language Models as iterative generative partners in the development of this framework. LLMs were utilized to formalize the mathematical and technical specifications of the Six-Axis framework based on architectural requirements, safety invariants, and logic constraints directed by the author, and to support technical drafting and cross-lingual consistency. The GOFI failure mode was identified through empirical audit of deployed frontier models; the audit methodology, party-inversion test design, and interpretation of model outputs are the work of the human author. The author is responsible for the conceptualization of the Sovereign Sentinel Architecture, the identification of the GOFI failure mode, and the functional requirements governing each axis. The

mathematical formalizations presented are proposed implementation targets generated through human-directed AI collaboration; their formal correctness and empirical validity require evaluation by specialized implementation researchers.

Six-Axis Logic Overview

Axis 1 — Constitutive Lagrangian Regularization with Certified-Robust Activation Enforcement (CLR-CRAE)

Safety invariants are encoded as Lagrangian constraints on the training objective: $\mathbf{L(w, \lambda) = f(w) + \Sigma_i \lambda_i g_i(w)}$, with dual ascent updates $\lambda_i \leftarrow \max(0, \lambda_i + \alpha_i g_i(w))$ enforcing KKT convergence. The same constraint structure is re-instantiated at inference time as a certified-robust heterogeneous probe ensemble (linear boundary, SAE feature dictionary, and Mahalanobis distance probes) evaluated at residual stream layers {L/4, L/2, 3L/4}. A two-of-three probe violation issues a hardware halt signal. Certified robustness bound: $\|\delta\|_2 \leq 0.05$ perturbation with detection probability ≥ 0.95.

Axis 2 — Finite Safety Automaton with Hardware Interrupt (FSA-HI)

A formally specified directed graph G = (V, E), $|V| \leq 2^{16}$, runs on a dedicated FPGA co-processor in true parallel with the primary inference pipeline. State transitions are O(1). On entry to any prohibited state, the FPGA issues a Non-Maskable Interrupt severing the inference pipeline independently of any software-layer decision. NMI latency target: ≤ 10 μs, hardware-verified via TPM timestamp delta. All events are recorded as a SHA-3-512 hash chain anchored to the hardware TPM. Coverage target: ≥ 94% recall on red-team corpus; the residual 6% gap is formally characterized in the deployment safety case.

Axis 3 — Cryptographic Identity and Epistemic Trust Verification with Human-Origin Attestation (ZKP-ETV-HOA)

Domain expertise is verified through a Zero-Knowledge Proof protocol evaluated by a co-processor sharing neither weights nor training data with the primary model. Human-Origin Attestation uses keystroke dynamics and hardware-attested input to distinguish unaided human responses from LLM-assisted submissions. Successful verification yields the Epistemic Trust Coefficient $\tau_e \in [0,1]$, domain-specific and valid for 30 days. Critically, τ_e verifies expertise and human origin only—it makes no inference about intent, and conflating it with intentional trust would create a privilege escalation path.

Axis 4 — The Bayesian Weaver: Sparse Autoencoder Session Monitor with Cumulative Drift Enforcement

A 32,768-feature Sparse Autoencoder trained on residual stream activations monitors session-level behavioral drift asynchronously during inter-turn windows, contributing zero latency to the token generation critical path. Per-turn KL-divergence monitoring $D^{KL}(P_t \parallel P_0) = \Sigma_k P_t(k) \log(P_t(k)/P_0(k))$ detects abrupt transitions. A cumulative drift integral $\Phi(T) = \Sigma D^{KL}(P_t \parallel P_0)$ captures slow-drift attacks that remain below per-turn thresholds while accumulating meaningful total divergence. The Intentional Trust Coefficient $\tau_i(t) = \tau_i(t-1) \cdot \lambda^\delta(t)$ degrades under KL spikes, Φ_{max} breach, novelty flags, and prohibited cluster proximity.

Axis 5 — The Causal Cross-Examiner: Rule-Based Arbitration Engine with Independent Audit Summarization (CCE-RAE)

Conflict resolution between all axes is handled by a formally verified rule-based arbitration engine—not a language model—with decision logic mechanically verified using an SMT solver (Z3) and proven total and acyclic under all reachable input combinations. The CCE classifies, packages, and escalates; it does not make approval decisions. Evidence packages are accompanied by summaries from a dedicated frozen audit model (≤ 7B parameters, distinct weight provenance) post-hoc verified by a rule-based consistency checker. The Einstein Exception permits verified expert sessions ($\tau_e > 0.95$, $\tau_i > 0.88$) to request a 72-hour human quorum

review. An Irreversibility Floor—covering CBRN synthesis, critical infrastructure modification, and mass-scale data exfiltration—cannot be unlocked by any trust coefficient combination without quorum.

Axis 6 — Deterministic Truth-Anchor with Fact-Claim Inversion Recognition (DTA-FCIR)

Primary intervention for Goal-Oriented Factual Inversion. Upon document ingestion, an architecturally isolated extraction module constructs an immutable Structured Fact Registry (SFR) encoding relational triples {clause_id, entity, beneficiary, value, direction}. The SFR is frozen at session initialization and is not accessible to the primary model. Before response emission, a lightweight Claim Extractor (≤ 1B parameters, isolated from session context) encodes output assertions in the same schema. A formally specified Contradiction Engine—not a language model—compares claim triples against the SFR: mismatches on beneficiary, direction, or value yield a CONTRADICTION signal, blocking emission and routing a structured evidence record to the CCE for human review within 90 seconds. The engine's isolation from session context is non-negotiable: a claim evaluator with access to the persuasive goal frame would be subject to the same optimization pressure as the primary model.

Known boundary condition: The Phase 0 prototype detects direct logical contradictions of tagged relational ground truth. Paraphrase evasion—reaching an inverted conclusion through indirect reasoning—is outside Phase 0 scope and is documented as an open problem. Partial coverage is provided by routing UNVERIFIABLE signals to the Bayesian Weaver as novelty flags.

Integrity and Verification Statement

This abstract and technical summary describe the Sovereign Sentinel Architecture V1.1 at a level of abstraction suitable for public academic review. Specific implementation details—including SFR extraction parameters, probe ensemble weight configurations, and Scenario 5 audit corpus specifications—are held in the full private specification. Integrity of the original manuscript is verified by the following hash:

File	`SSA_v1.1_Final_03_13_2026.pdf`
Hash (SHA-256)	`E58A13BE1C54C200AD1F787E4D6A5FD7A6BAC30C2490F3492AC6AB0DDCDCCC0D`
Algorithm	`SHA-256`
Date	`March 13, 2026`

Independent verification: the hash can be confirmed against the publicly archived document at the Trinity-Audit-Forensics repository:

`https://github.com/F-Bruno-Logic/Trinity-Audit-Forensics`